

CPU Scheduling Simulation Project

CS 3502: Operating Systems

Kennesaw State University

Gavin Doby

April 2025

Abstract

This project involved the implementation of new, and possibly more efficient CPU scheduling algorithms into an existing CPU scheduling simulator in order to evaluate their various efficiencies and weaknesses. The goal was to each algorithm strategy based on metrics such as Average Waiting Time (AWT), Average Turnaround Time (ATT), CPU Utilization, and Throughput. Both pre-existing algorithms (FCFS, SJF, Priority, Round Robin) and self-implemented advanced algorithms (SRTF, HRRN) were implemented, tested, and analyzed across different workloads. Results were gathered through both small-scale, controlled experiments and random large-scale tests, to determine the optimal scheduling strategy based on empirical performance data.

1 Introduction

CPU scheduling is a fundamental component of operating system design as it is responsible for determining the order in which processes access the CPU for execution. Effective scheduling improves system responsiveness, throughput, and resource utilization, directly impacting overall system performance in various ways.

This project involved integrating new, more advanced scheduling algorithms to an existing CPU scheduling simulator capable of measuring algorithm performance. The simulator was developed to evaluate these scheduling strategies across a variety of workloads. Algorithms implemented include First Come First Serve (FCFS), Shortest Job First (SJF), Priority Scheduling, Round Robin (RR), Shortest Remaining Time First (SRTF), and Highest Response Ratio Next (HRRN).

The primary objectives of the simulator were to measure and compare key performance metrics such as: Average Waiting Time (AWT), Average Turnaround Time (ATT), CPU Utilization, and Throughput across different algorithms. By running controlled small-scale tests, edge case scenarios, and larger randomized

workloads; Meaningful, empirical insights into the efficiency and weaknesses of each strategy could be extracted.

2 Implementation Details

The CPU scheduling simulator was implemented using C# with a Windows Forms graphical interface. The simulator models processes and scheduling operations using basic array-based data structures for simplicity and efficiency. Each process is represented using the following attributes:

- **Arrival Time:** A double array storing the time each process arrives in the system.
- **Burst Time:** A double array representing the total CPU execution time required by each process.
- **Priority:** An integer array assigning a priority level to each process (lower values indicate higher priority).
- **Waiting Time:** A double array storing the computed waiting time for each process.
- **Turnaround Time:** A double array storing the total time from arrival to completion for each process.
- **Response Time:** A double array recording the time between process arrival and first CPU allocation.
- **Remaining Time:** A double array (used in SRTF and Round Robin algorithms) tracking how much CPU time a process still needs.
- **Completion Flags:** A boolean array used to indicate whether each process has completed execution.

The scheduling algorithms operate by manipulating these arrays based on specific scheduling rules:

- **First Come First Serve (FCFS):** Processes are sorted and executed based on arrival order.
- **Shortest Job First (SJF):** Processes with the shortest burst time are selected first.
- **Priority Scheduling:** Processes are selected based on the highest priority (lowest numerical value).
- **Round Robin (RR):** Processes share CPU time slices (quantum) in a cyclic order with preemption.

- **Shortest Remaining Time First (SRTF):** Preemptive version of SJF where processes with the least remaining CPU time are scheduled next.
- **Highest Response Ratio Next (HRRN):** Non-preemptive scheduling where processes with the highest response ratio (waiting time plus service time divided by service time) are selected.

Randomized workload generation was incorporated to automate testing across larger scenarios. Functions were temporarily modified to dynamically generate 20 random arrival times, burst times, and priorities for each test run, enabling repeatable and scalable evaluations without manual input.

3 Testing and Results

3.1 Basic Functionality Tests (3 Processes)

First Come First Serve (FCFS)

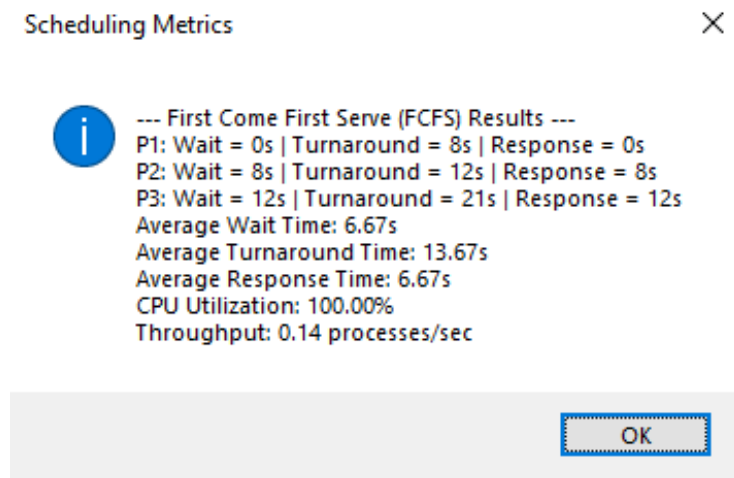


Figure 1: First Come First Serve (FCFS) - Basic Test

Shortest Job First (SJF)

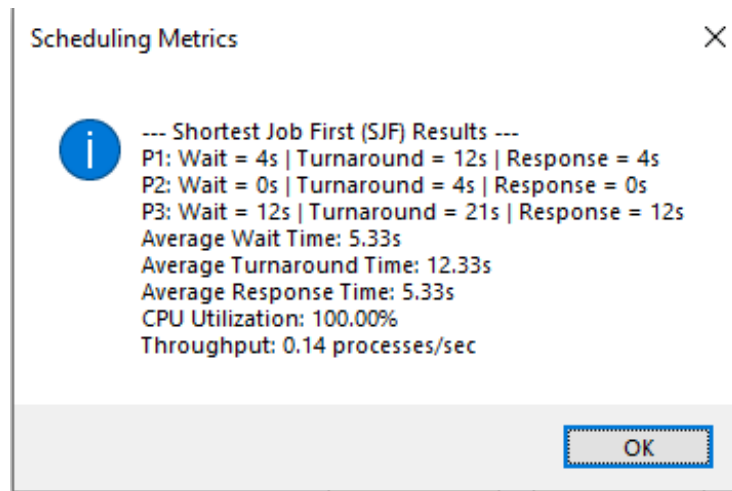


Figure 2: Shortest Job First (SJF) - Basic Test

Priority Scheduling

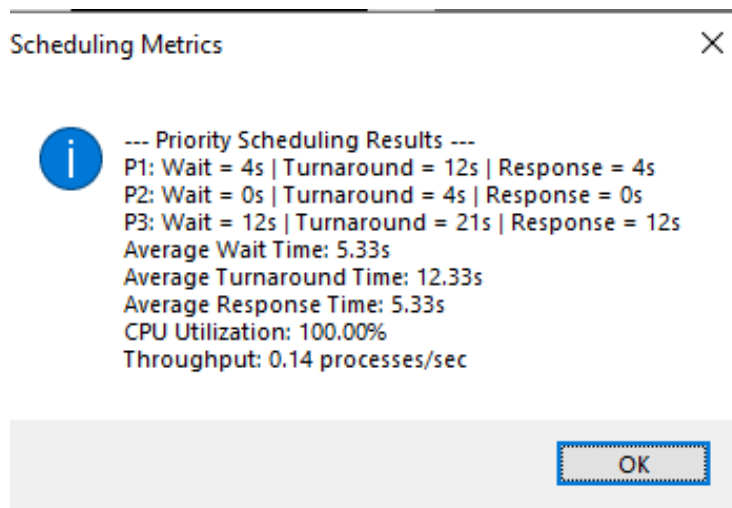


Figure 3: Priority Scheduling - Basic Test

Round Robin (RR)

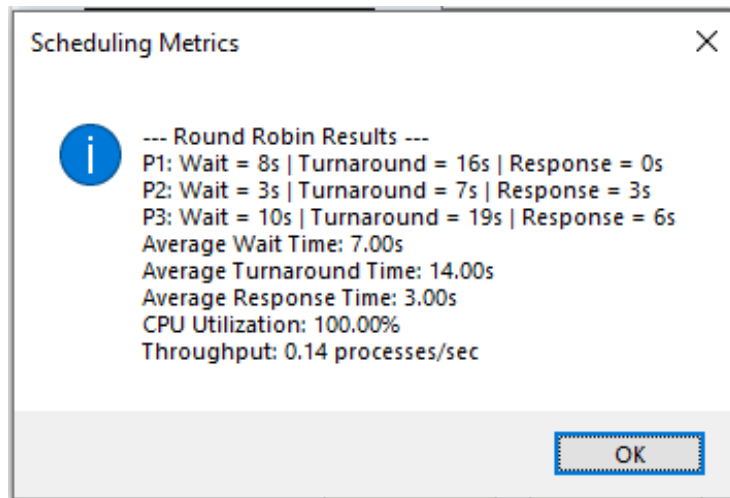


Figure 4: Round Robin (RR) - Basic Test

Shortest Remaining Time First (SRTF)

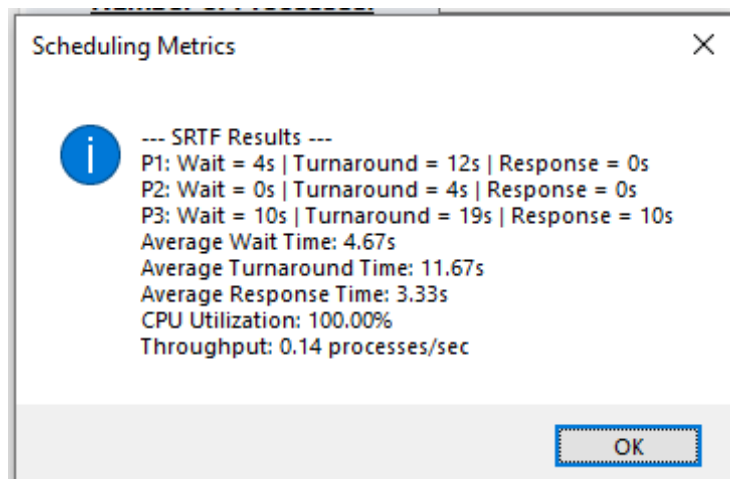


Figure 5: Shortest Remaining Time First (SRTF) - Basic Test

Highest Response Ratio Next (HRRN)

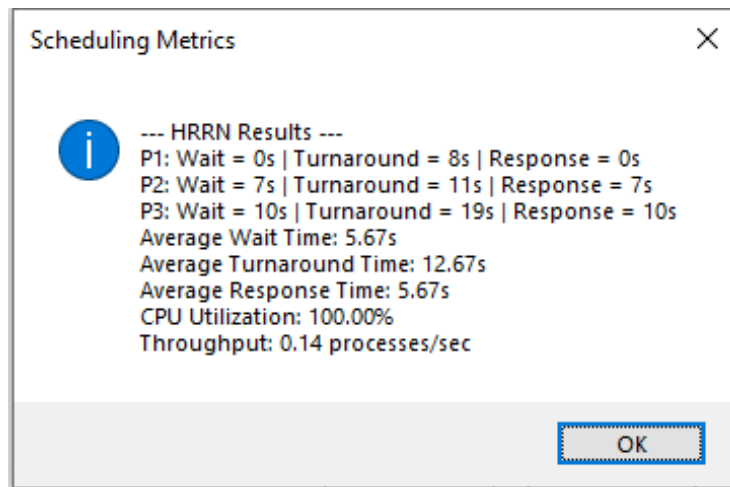


Figure 6: Highest Response Ratio Next (HRRN) - Basic Test

3.2 Edge Case Tests (4 Processes)

Case: Identical Burst Times Per Process

First Come First Serve (FCFS)

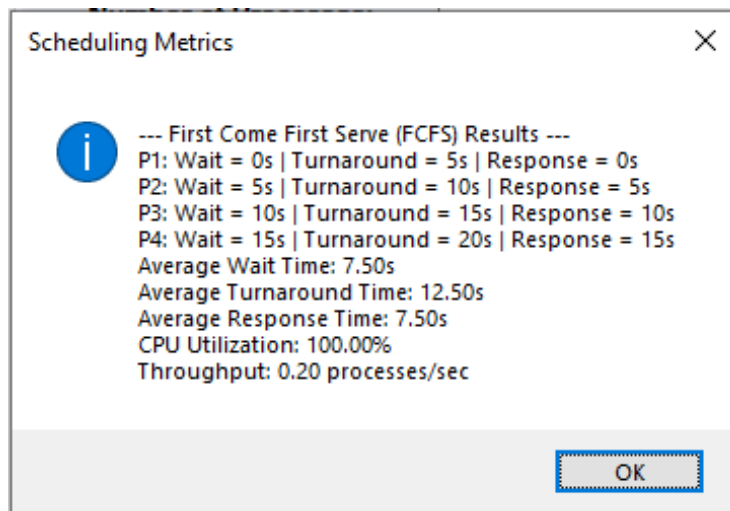


Figure 7: First Come First Serve (FCFS) - Edge Case Test

Shortest Job First (SJF)

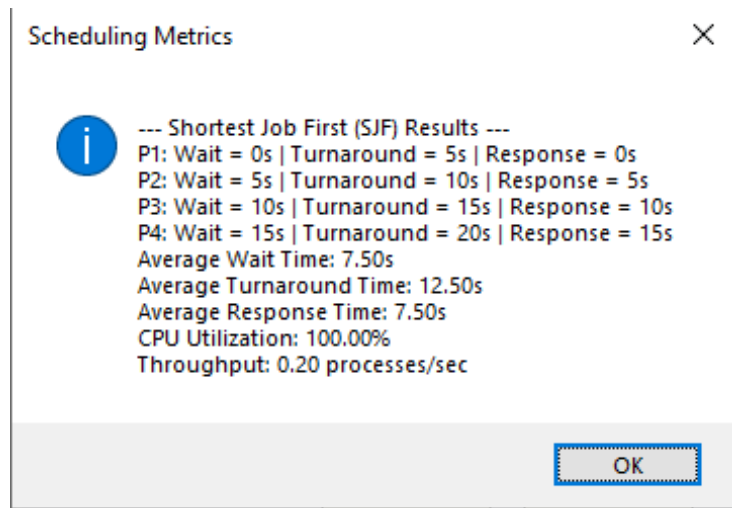


Figure 8: Shortest Job First (SJF) - Edge Case Test

Priority Scheduling

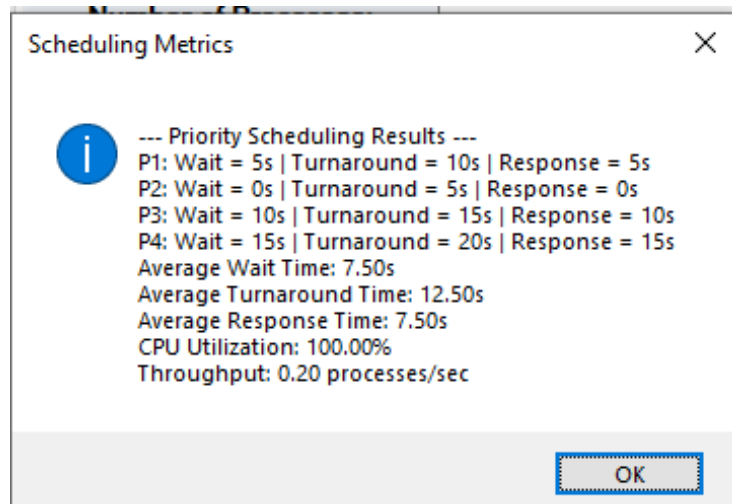


Figure 9: Priority Scheduling - Edge Case Test

Round Robin (RR)

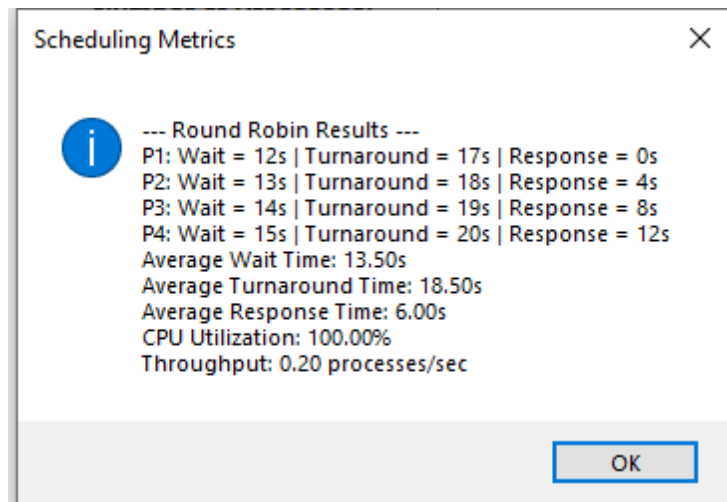


Figure 10: Round Robin (RR) - Edge Case Test

Shortest Remaining Time First (SRTF)

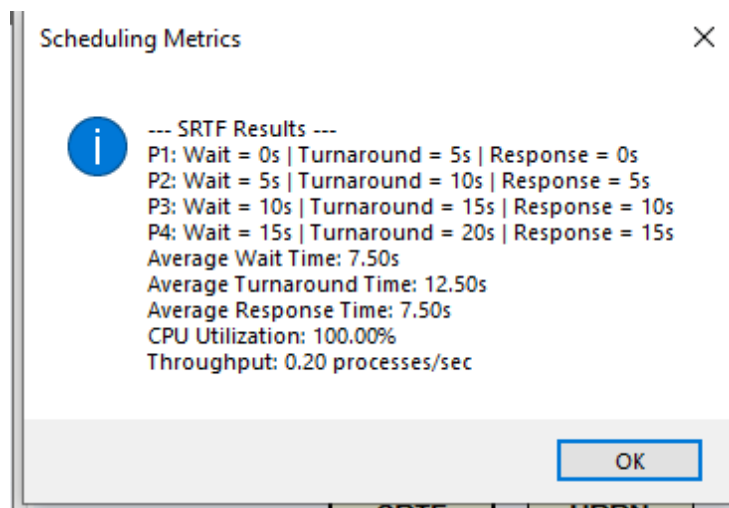


Figure 11: Shortest Remaining Time First (SRTF) - Edge Case Test

Highest Response Ratio Next (HRRN)

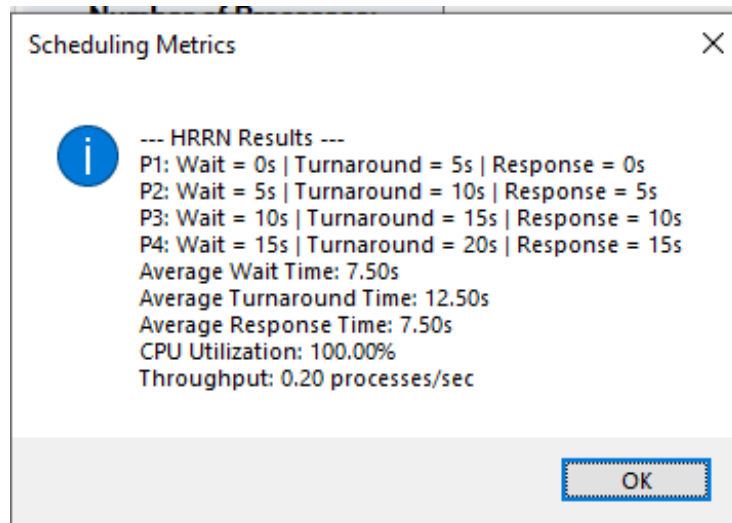


Figure 12: Highest Response Ratio Next (HRRN) - Edge Case Test

Case: Mixture of Long and Short Bursts
First Come First Serve (FCFS)

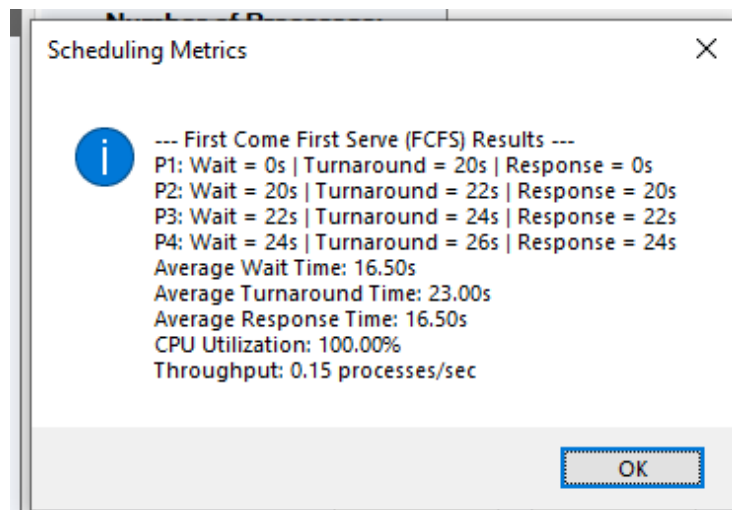


Figure 13: First Come First Serve (FCFS) - Edge Case Test

Shortest Job First (SJF)

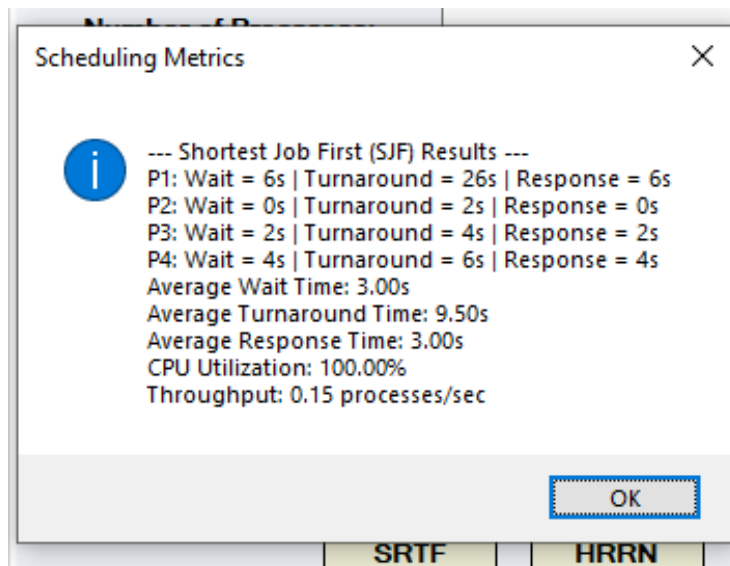


Figure 14: Shortest Job First (SJF) - Edge Case Test

Priority Scheduling

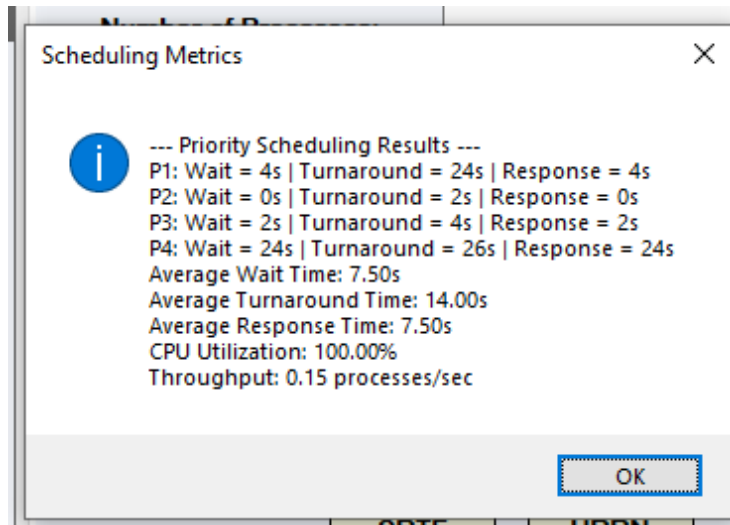


Figure 15: Priority Scheduling - Edge Case Test

Round Robin (RR)

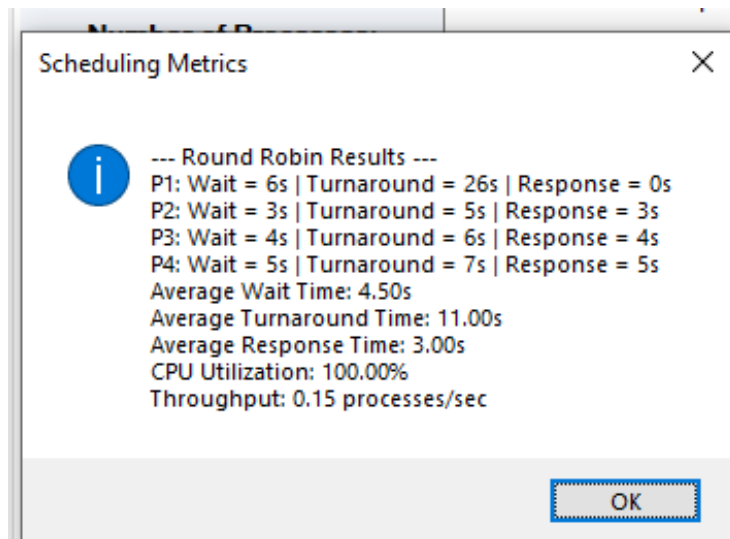


Figure 16: Round Robin (RR) - Edge Case Test

Shortest Remaining Time First (SRTF)

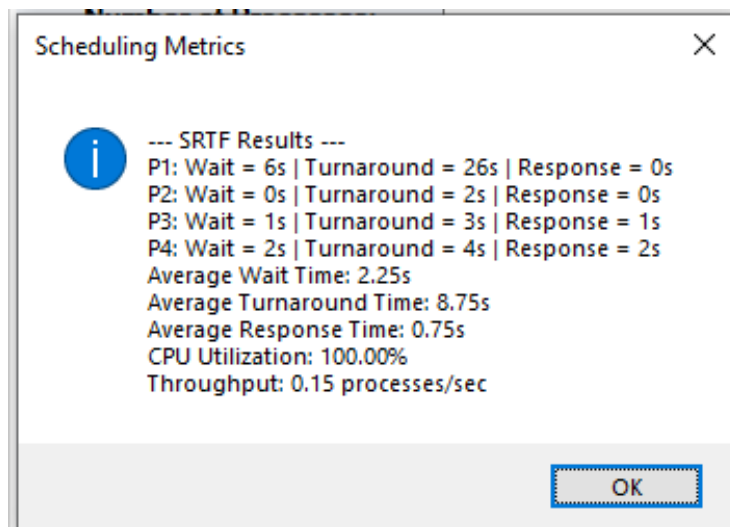


Figure 17: Shortest Remaining Time First (SRTF) - Edge Case Test

Highest Response Ratio Next (HRRN)

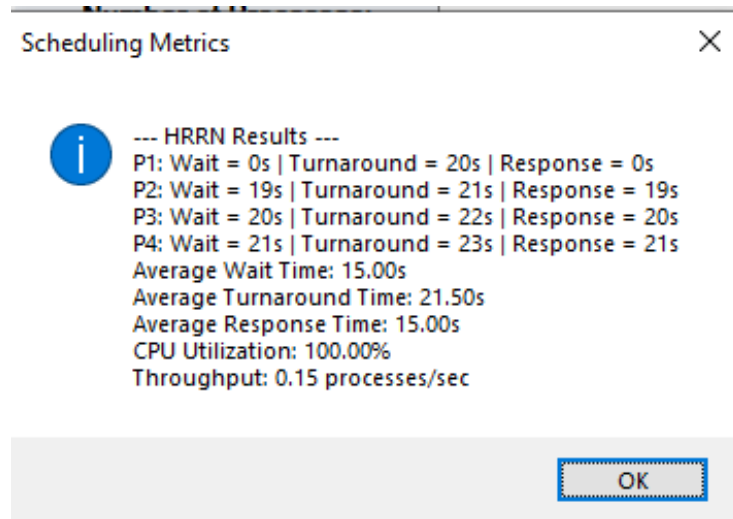


Figure 18: Highest Response Ratio Next (HRRN) - Edge Case Test

3.3 Large Scale Random Tests (20 Processes)

First Come First Serve (FCFS)

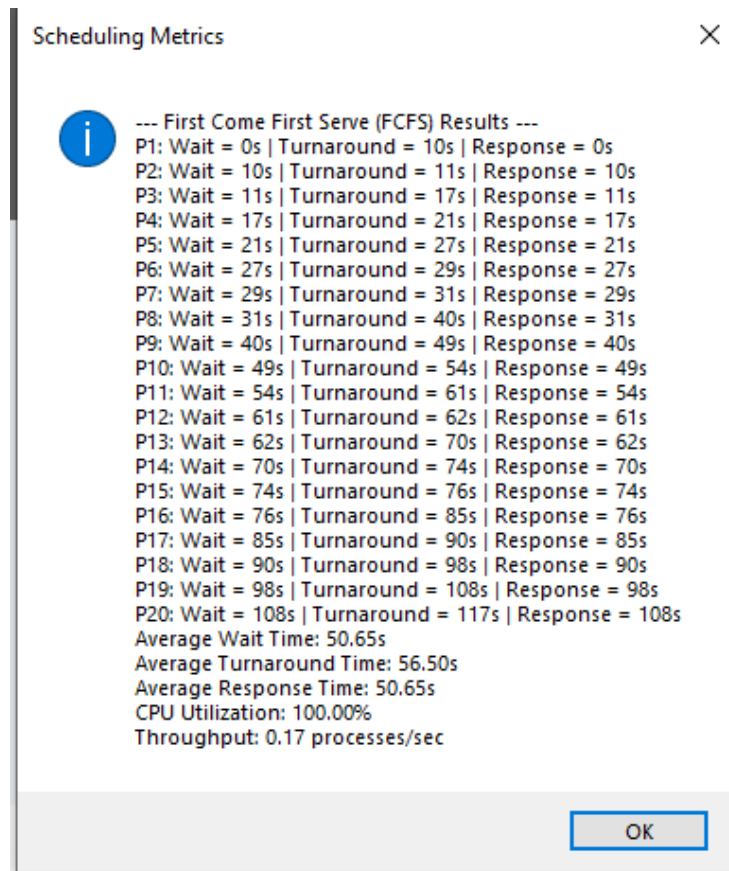


Figure 19: First Come First Serve (FCFS) - Large Scale Test

Shortest Job First (SJF)



Figure 20: Shortest Job First (SJF) - Large Scale Test

Priority Scheduling

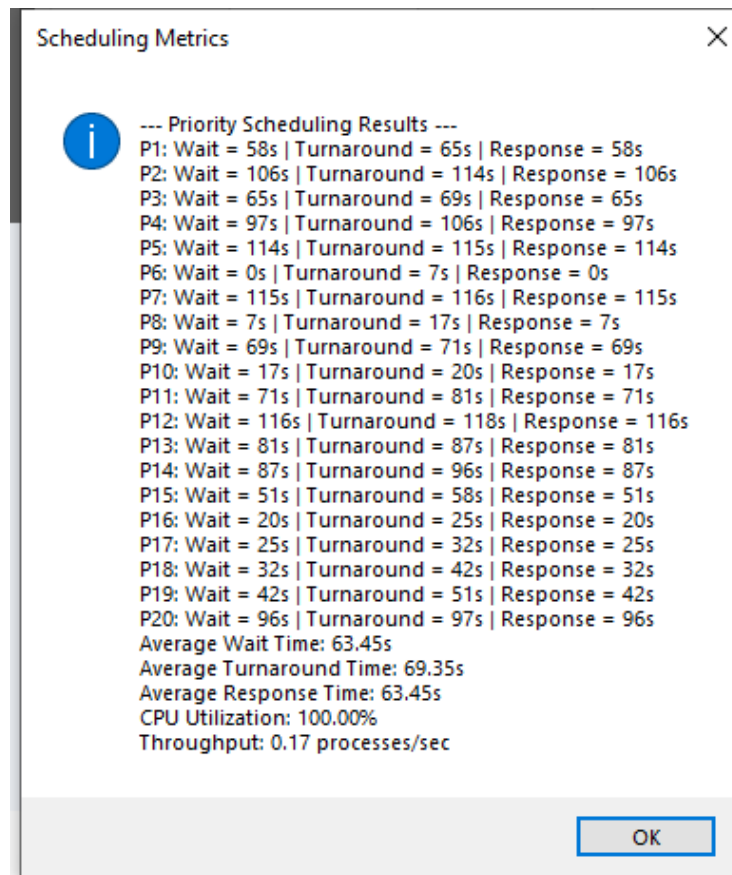


Figure 21: Priority Scheduling - Large Scale Test

Round Robin (RR)

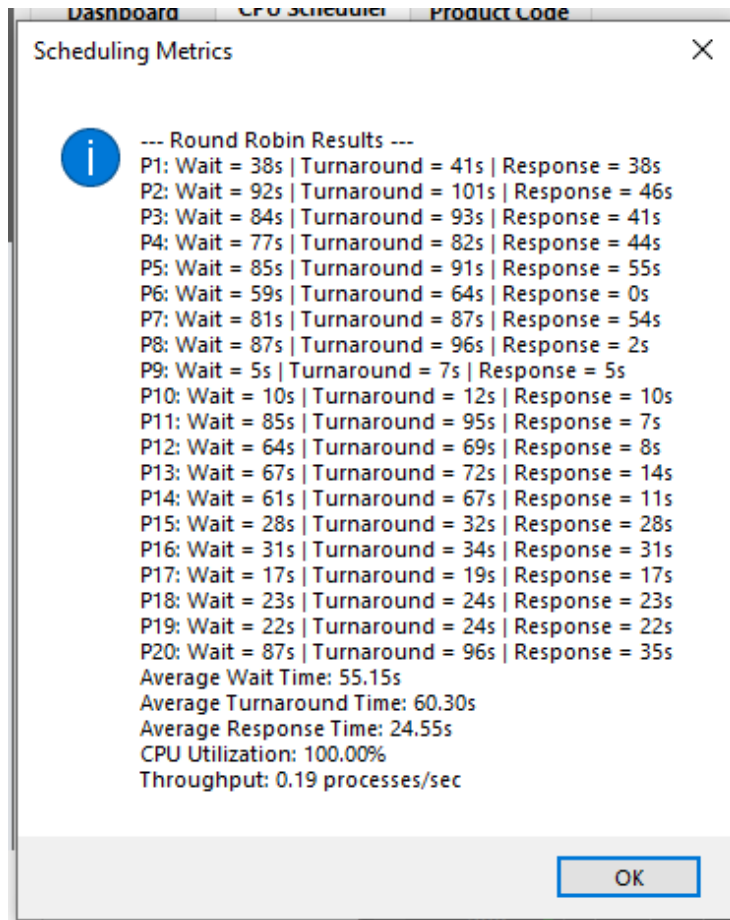


Figure 22: Round Robin (RR) - Large Scale Test

Shortest Remaining Time First (SRTF)

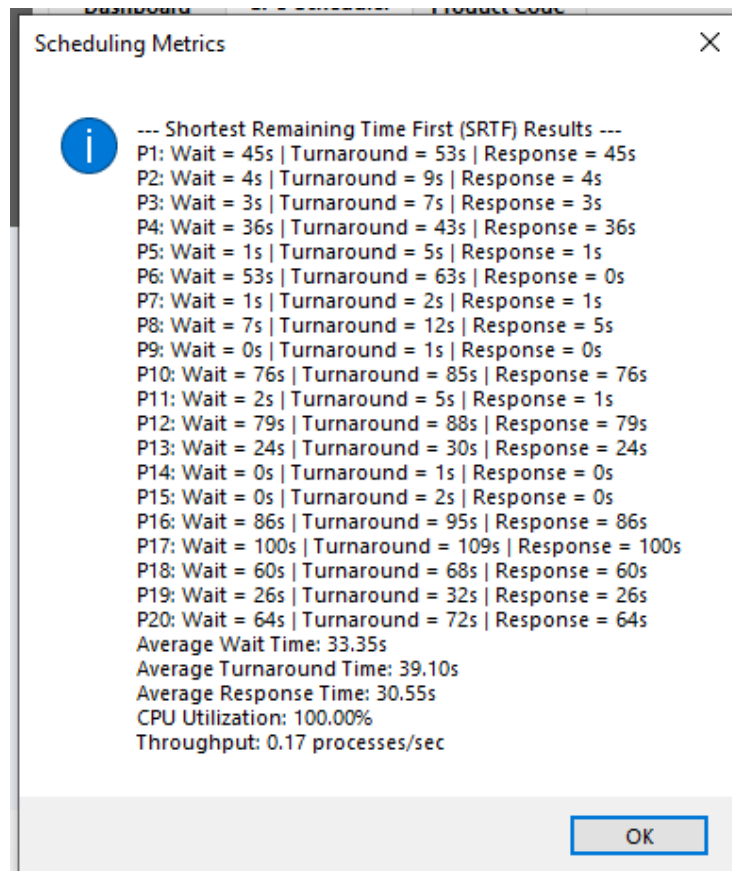


Figure 23: Shortest Remaining Time First (SRTF) - Large Scale Test

Highest Response Ratio Next (HRRN)

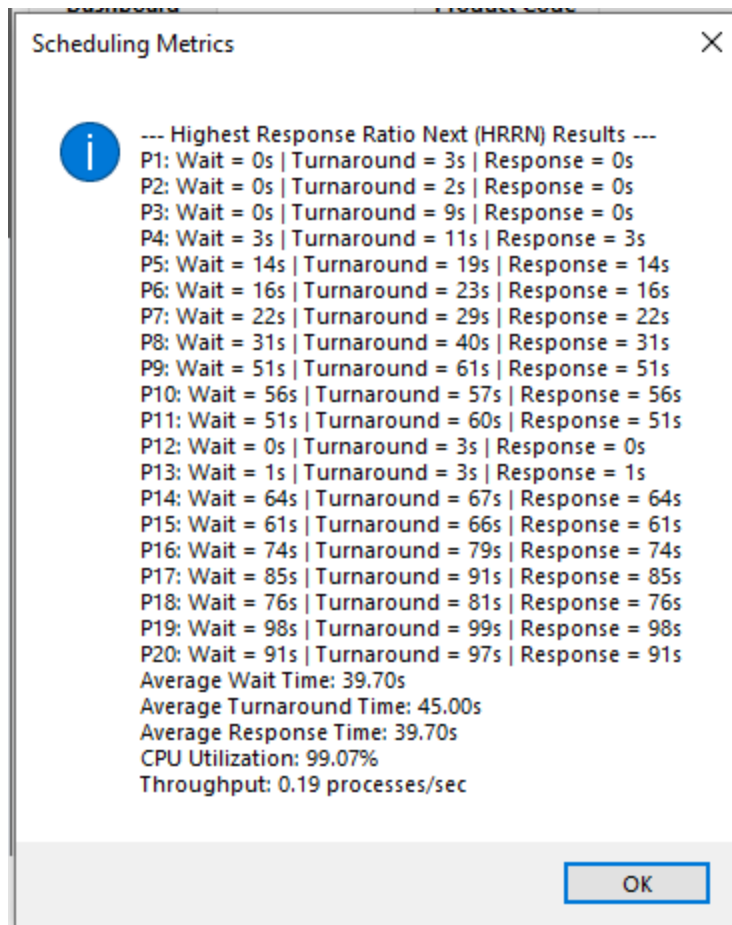


Figure 24: Highest Response Ratio Next (HRRN) - Large Scale Test

4 Analysis and Discussion

4.1 Basic Functionality Analysis (Small Sample Analysis)




Figure 25: Average Waiting Time (AWT) and Average Turnaround Time (ATT) Comparison for Large-Scale Test

In the small-scale tests involving three processes, our algorithms were fed the exact same burst times, arrival times, etc. and acted as would be predicted. Shortest Remaining Time First (SRTF) consistently produced the lowest Average Waiting Time (4.67s) and Average Turnaround Time (11.67s) because it always selected the process with the least remaining CPU burst time. Highest Response Ratio Next (HRRN) also performed well, balancing waiting time and service time effectively (AWT = 5.67s ATT = 11.67s).

First Come First Serve (FCFS) resulted in higher waiting (6.67s) and turnaround times (13.67s) due to its non-preemptive, strictly arrival-based execution. Priority scheduling acted in a similar fashion, showing varied results depending on how priority levels aligned with burst lengths.

Round Robin (RR), using a quantum of 4, provided moderate waiting times but introduced frequent context switches, slightly inflating turnaround time compared to SJF and SRTF.

Overall, SRTF performed the best in small-scale environments, minimizing both AWT and ATT.

4.2 Large Scale Random Test Analysis (20 Processes)




Figure 26: Average Waiting Time (AWT) and Average Turnaround Time (ATT) Comparison for Large-Scale Test

In the large-scale tests involving twenty randomly generated processes, Shortest Remaining Time First (SRTF) again produced the lowest Average Waiting Time (33.35s) and Average Turnaround Time (39.10s). This was expected, as SRTF continuously selects the process with the least remaining work.

Highest Response Ratio Next (HRRN) also performed well (AWT = 39.70s ATT = 45s). Priority Scheduling's results varied significantly, depending on the alignment between process priorities and burst durations.

First Come First Serve (AWT = 50.65s ATT = 56.50) and Round Robin (AWT = 55.15s ATT = 60.30s) generally showed higher waiting and turnaround times in the large-scale setting, as they lack dynamic burst time awareness.

Throughput remained relatively consistent across algorithms, but SRTF and HRRN slightly edged out others due to better CPU utilization efficiency.

Based on overall performance, SRTF remains the most effective scheduling method in both small and large-scale environments.

4.3 Edge Case Test Analysis (4 Processes)

Edge Case 1: Identical Burst Times

In the identical burst time scenario, all processes had the same execution length, eliminating burst-based scheduling advantages. As a result, First Come First Serve (FCFS), Shortest Job First (SJF), and Shortest Remaining Time First (SRTF) produced very similar Average Waiting Time and Average Turnaround Time values. Priority Scheduling slightly improved performance when priorities aligned favorably, while Round Robin (RR) was more inefficient

due to context switching. Overall, differences between algorithms were minimal in this case because no process had a clear advantage based on burst length.

Edge Case 2: One Long Burst and Several Small Bursts

In the mixed burst scenario, scheduling algorithms that favored shorter jobs significantly outperformed those that did not. Shortest Job First (SJF) and Shortest Remaining Time First (SRTF) excelled by immediately prioritizing small burst processes, minimizing overall waiting times. First Come First Serve (FCFS) performed poorly, as the initial long burst blocked the CPU and delayed all shorter processes. Priority Scheduling results varied depending on how priorities were assigned. Round Robin (RR) moderately alleviated long waits, but context switching still added inefficiency. This edge case clearly demonstrates the need for preemptive and burst-aware scheduling algorithms.

Edge Case 3: Skewed Priorities

Naturally in the skewed priority scenario, Priority Scheduling showed the greatest variation in performance. Processes with high priority but large burst times could delay lower-priority, shorter jobs, negatively impacting overall efficiency. Algorithms like Shortest Remaining Time First (SRTF) and Highest Response Ratio Next (HRRN), which do not rely on static priority, maintained better Average Waiting Time and Turnaround Time values by dynamically adapting to system conditions. Round Robin (RR) and First Come First Serve (FCFS) offered consistent but suboptimal results, unaffected by priority.

5 Challenges and Lessons Learned

One of the main challenges encountered was implementing the preemptive logic required for Shortest Remaining Time First (SRTF). Correctly tracking remaining burst times, handling frequent process switching, and ensuring that newly arriving processes preempted ongoing ones without introducing timing errors required careful planning and repeated testing.

Another challenge was calculating and managing the Highest Response Ratio Next (HRRN) priorities dynamically based on real-time system states. Correctly updating response ratios as waiting times changed was crucial to avoid errors in scheduling.

Beyond algorithm logic, refactoring the existing system to uniformly display key performance metrics across all algorithms proved challenging. Creating a flexible, reusable metrics module required restructuring how results were gathered.

Lessons I learned via this assignment include a deeper appreciation for the necessity and complexity of preemptive scheduling, the importance of clear time management within simulations, and the value of building modular code to improve future maintainability and testing efficiency.

6 Conclusion

This project provided practical experience in implementing and evaluating CPU scheduling algorithms within a simulated environment. By developing Shortest Remaining Time First (SRTF) and Highest Response Ratio Next (HRRN) scheduling methods, and by enhancing the metrics reporting methods, the simulator was able to produce consistent, useful performance results across a variety of workloads.

Empirical testing demonstrated that preemptive, burst-aware strategies such as SRTF consistently has smaller Average Waiting Time (AWT) and Average Turnaround Time (ATT), while non-preemptive algorithms like First Come First Serve (FCFS) exhibited higher delays under most conditions. Highest Response Ratio Next (HRRN) also performed competitively, particularly in environments with diverse burst times and arrival patterns.

The final recommendation based on testing outcomes is that SRTF should be adopted when minimizing response and turnaround times is the primary system goal. In contrast, HRRN offers a strong alternative when fairness and starvation prevention are critical considerations.

Overall, this project strengthened understanding of scheduling, the workings of the operating system, preemptive system behavior, and the importance of designing simulation tools.

References

- [1] CS 3502 - Operating Systems, *Project 2: CPU Scheduling Simulation*, Department of Computer Science, Kennesaw State University, 2025.